

PCs that are ready for local AI

A free OS that makes local AI effortless — and the machines that run it out of the box.

One Pager

2026-06-29

Phase 0 · prototype

THE PROBLEM

Local AI is finally capable and the open models are good enough — but the setup is miserable: CUDA/ROCm versions, Python & venv breakage, "which model fits my VRAM?", exposing services safely. Most people bounce off it.

THE PRODUCT

CloudlessOS — a Linux distro that installs and runs AI apps (ComfyUI, vLLM, Open WebUI...) as clean GPU containers from a friendly web UI. Plus **Cloudless PCs** that ship with it pre-tuned.

HOW IT WORKS — THREE REAL LAYERS UNDER A THIN SHELL

1 · Orchestrator installs / runs / updates AI apps and manages models & GPUs · **2 · Hardware enablement** handles drivers, CUDA/ROCm, VRAM detection & "fits your machine" guidance · **3 · Curated catalog** of vetted recipes so one-click installs are reliable.

THE TWO STARTING MACHINES

Cloudless Spark

NVIDIA GB10 · 128 GB · CUDA

No porting, ~\$3.5–4k, reseller path.

Start here.

Cloudless Mini

AMD Strix Halo · 128 GB · x86

~\$2k, own-brand path; porting off-CUDA needed.

THE BUSINESS MODEL

Free OS = the funnel. The machine = the revenue. You never sell software — you sell a box where it already works.

A free download earns trust the way Pop!_OS sells System76 laptops; the configured machine is the easy yes that follows.

ROADMAP — SOFTWARE FIRST, HARDWARE LAST

CURRENT

Phase 0 — Orchestrator

The magic moment: click an app → it installs into a GPU container, opens, uninstalls cleanly.

NEXT

Phase 1 — OS image

Bootable CloudlessOS + kiosk shell. Publish the free download — the funnel goes live.

LATER

Phase 2 — Hardware

Reference machines pre-tuned for local AI, once the software is the reason to buy.

GO TO MARKET — CAPITAL-LIGHT, FUNDED BY CUSTOMERS

0 Buy one box

1 Ship the free OS

2 Build-to-order (pay first → no inventory)

3 Crowdfund the first private-label run

4 ODM white-label with traction